

# Mufei Li

mufeili1996@gmail.com

Homepage: <https://mufeili.github.io>

## Education

- |   |   |
|---|---|
| <ul style="list-style-type: none"><li>• <b>Georgia Institute of Technology</b><br/><i>Ph.D., Machine Learning</i></li><li>• <b>New York University Shanghai</b><br/><i>B.S., Honors Mathematics</i></li></ul> | <p>Georgia, United States<br/><i>Aug 2023 - May 2027 (expected)</i></p> <p>Shanghai, China<br/><i>Sep 2014 - May 2018</i></p> |
|---|---|

## Professional Experience

- |   |   |
|---|---|
| <ul style="list-style-type: none"><li>• <b>Research Scientist Intern</b><br/><i>Meta Ranking and Foundational AI</i></li><li>• <b>Software Development Engineer II</b><br/><i>Amazon Web Services Shanghai AI Lab</i></li><li>• <b>Software Development Engineer I</b><br/><i>Amazon Web Services Shanghai AI Lab</i></li><li>• <b>Research Assistant</b><br/><i>New York University Shanghai</i></li></ul> | <p>California, United States<br/><i>May 2025 - Nov 2025</i></p> <p>Shanghai, China<br/><i>Jul 2020 - Jun 2023</i></p> <p>Shanghai, China<br/><i>Feb 2019 - Jul 2020</i></p> <p>Shanghai, China<br/><i>Jun 2018 - Feb 2019</i></p> |
|---|---|

## Selected Publications

- Mufei Li, Dongqi Fu, Limei Wang, Si Zhang, Hanqing Zeng, Kaan Sancak, Ruizhong Qiu, Haoyu Peter Wang, Xiaoxin He, Xavier Bresson, Yinglong Xia, Chonglin Sun, Pan Li, **Haystack Engineering: Context Engineering for Heterogeneous and Agentic Long-Context Evaluation**, NeurIPS 2025 workshop on Evaluating the Evolving LLM Lifecycle.
- Xinjie Shen, Mufei Li, Pan Li, **Measuring Physical-World Privacy Awareness of Large Language Models: An Evaluation Benchmark**, NeurIPS 2025 workshop on Bridging Language, Agent, and World Models.
- Haoyu Wang, Peihao Wang, Mufei Li, Shikun Liu, Siqi Miao, Zhangyang Wang, Pan Li, **Graph-KV: Breaking Sequence via Injecting Structural Biases into Large Language Models**, Advances in Neural Information Processing Systems (NeurIPS), 2025.
- Rongzhe Wei, Mufei Li, Mohsen Ghassemi, Eleonora Kreacić, Yifan Li, Xiang Yue, Bo Li, Vamsi K. Potluru, Pan Li, Eli Chien, **Underestimated Privacy Risks for Minority Populations in Large Language Model Unlearning**, International Conference on Machine Learning (ICML), 2025.
- Mufei Li\*, Siqi Miao\*, Pan Li, **Simple is Effective: The Roles of Graphs and Large Language Models in Knowledge-Graph-Based Retrieval-Augmented Generation**, International Conference on Learning Representations (ICLR), 2025 (\*Equal contribution); also selected for **oral presentation** at ICLR 2025 workshop on Foundation Models in the Wild.

- Mufei Li, Viraj Shitole, Eli Chien, Changhai Man, Zhaodong Wang, Srinivas Sridharan, Ying Zhang, Tushar Krishna, Pan Li, **LayerDAG: A Layerwise Autoregressive Diffusion Model for Directed Acyclic Graph Generation**, International Conference on Learning Representations (ICLR) **Spotlight (top 5.1%)**, 2025; also selected for **spotlight presentation** at ICLR 2025 workshop on “Will Synthetic Data Finally Solve the Data Access Problem?”.
- Mufei Li, Jinjing Zhou, Jiajing Hu, Wenzuan Fan, Yangkang Zhang, Yixin Gu, George Karypis, **DGL-LifeSci: An Open-Source Toolkit for Deep Learning on Graphs in Life Science**, ACS Omega, 2021. **700 stars on GitHub**.
- Minjie Wang, Da Zheng, Zihao Ye, Quan Gan, Mufei Li, Xiang Song, Jinjing Zhou, Chao Ma, Lingfan Yu, Yu Gai, Tianjun Xiao, Tong He, George Karypis, Jinyang Li, Zheng Zhang, **Deep Graph Library: A Graph-Centric, Highly-Performant Package for Graph Neural Networks**, arXiv, 2019. **14K stars on GitHub**.

## Honors and Awards

- Notable Reviewer for ICLR 2025
- Top Reviewer for NeurIPS 2024
- Georgia Tech ECE Fellowship

## Peer Review Service

- Conference on Neural Information Processing Systems (NeurIPS), 2024 - 2025
- International Conference on Learning Representations (ICLR), 2025 - 2026
- International Conference on Machine Learning (ICML), 2025
- IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI), 2024
- Transactions on Machine Learning Research (TMLR), 2025
- ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD), 2024
- Learning on Graphs Conference (LoG), 2024
- Asian Conference on Machine Learning (ACML), 2024

## Teaching Assistant

- ECE 6254 Statistical Machine Learning